

Topic_Modeling

Group 10: Bill Gao, Jiun Lee, Priam Vyas, Danya Zhang

2022-11-16

```
##Import Data from kaggle
```

```
##Import Data
```

```
imdb <- read_csv ("/Users/jiunlee/MSSP22/Fidelity/IMDB Dataset.csv",  
                  col_names = TRUE,  
                  show_col_types = FALSE)
```

```
##Remove sentiment column since we don't use it.
```

```
imdb <- imdb %>% select(-sentiment)
```

```
##Remove duplicates of review, now we have 49582 observations.
```

```
imdb <- unique(imdb)
```

```
##Sample down for 100 reviews
```

```
set.seed(456)
```

```
review_index <- 1:dim(imdb)[1]
```

```
text_df <- cbind(review_index,imdb)
```

```
text_df <- text_df %>% slice_sample(n = 100, replace = FALSE)
```

```
rm(review_index)
```

```
##After a thorough cleaning, we now have a random sample of 100 observations data frame
```

```
##Cleaning: Tokenizing, Removing stopwords,tf-idf
```

```
##The imdb dataset has a lot of stopwords and meaningless words. #We will remove stopwords and words un
```

```
library(stringr)
```

```
library(tidytext)
```

```
## Tokenizing, count the number of words within each review.
```

```
token <- text_df %>%
```

```
  unnest_tokens(word, review) %>%
```

```
  count(review_index, word, sort=TRUE) %>%
```

```
  rename(count=n)
```

```
##There's a lot of stop words. Let's remove them.
```

```
##Create a stop word vector
```

```
stop <- unlist(stop_words[,1])
```

```
##Drop the attribute
```

```
stop <- StripAttr(stop)
```

```
##Restore tokens Data set
```

```
check <- token
```

```

##Check stop word lists again
remove <- check$word %in% stop
##To make it easier to see, create a data frame
d <- cbind(token,remove)
##Create an index of words(not stopwords)
f <- which(d$remove == FALSE)
##Clean tokens that has no stopwords
clean_token <- d %>% slice(f) %>% select(-remove)

##Let's subset the Clean-Token

##Vector that has meaningless words
strings <- c("br","movie","film", "scene", "character","story","bit","lot","bad","act","hard","awful",")

##Detect numbers of rows that has meaningless words
meaningless <- str_detect(clean_token$word, paste(strings, collapse = "|"))

##Detect numbers of meaningful rows
meaningful <- which(meaningless==F)

##Subset: tokens without meaningless
clean_token <- clean_token %>% slice(meaningful)

##Remove redundant data and values
rm(d,check,f,meaningless,meaningful,strings,stop,remove,token)

##Now we have our new clean_token data with only 3776 observations

```

##Create Latent Dirichelet Allocation model

```

library(topicmodels)

##Convert sample token tibble to document term matrix for LDA
clean_token_dmat <- clean_token %>%
  cast_dtm(review_index, word, count)

##Select k= 6 because we have 6 general film genres
imdb_lda <- LDA(clean_token_dmat, k = 6, control = list(seed = 1234))
imdb_lda #A LDA_VEM topic model with 6 topics.

```

A LDA_VEM topic model with 6 topics.

```
imdb_topics <- tidy(imdb_lda, matrix = "beta")
```

##Now we create our plots

```

library(ggplot2)
##Get top used terms and arrange them
imdb_top_terms <- imdb_topics %>%
  group_by(topic) %>%
  slice_max(beta, n = 6) %>%
  ungroup() %>%

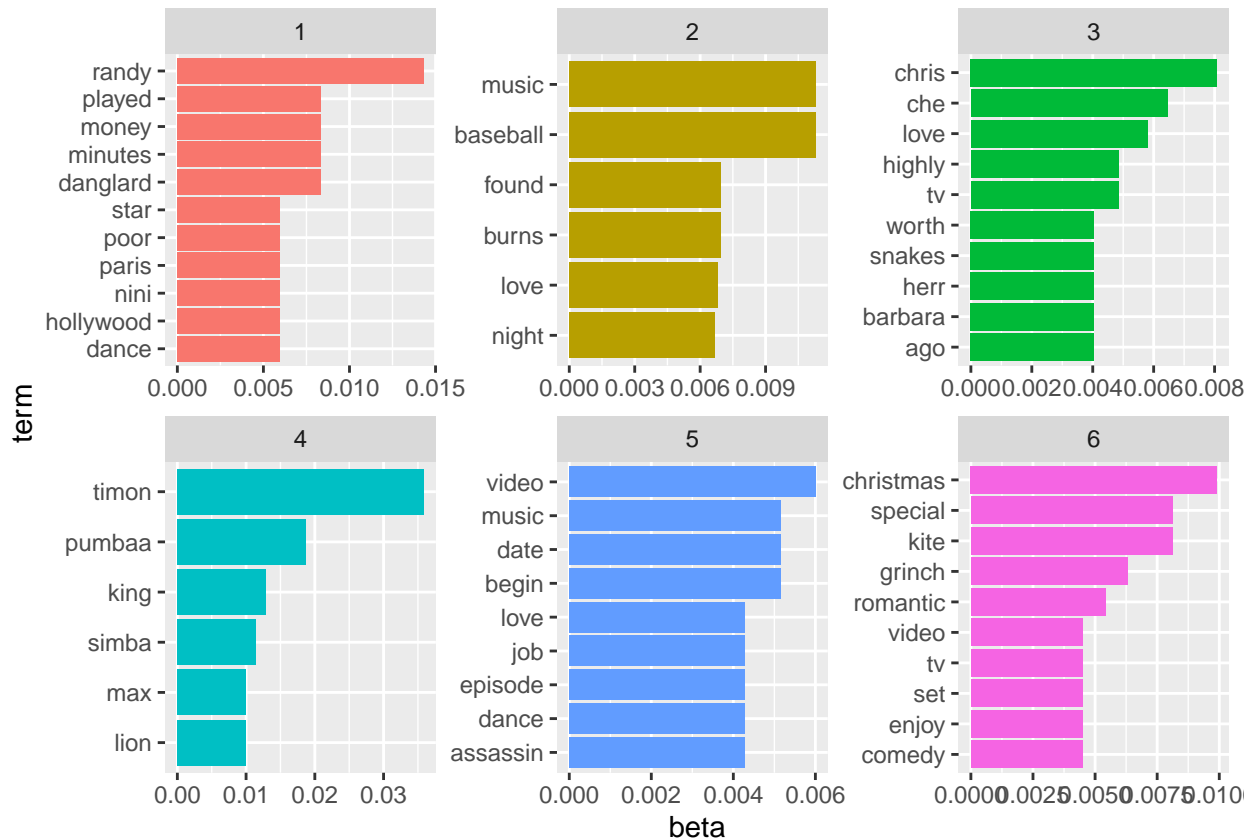
```

```

arrange(topic, -beta)

##Create the plot
imdb_top_terms %>%
  mutate(term = reorder_within(term, beta, topic)) %>%
  ggplot(aes(beta, term, fill = factor(topic))) +
  geom_col(show.legend = FALSE) +
  facet_wrap(~ topic, scales = "free") +
  scale_y_reordered()

```



##The plot above shows the top words for each topic in our LDA.
 ##We've split up our LDA into 6 genres, which represents the number of topics we have.

##Now let's look at Document-topic probabilities

```

##Gamma: per-document-per-topic probabilities
imdb_documents <- tidy(imdb_lda, matrix = "gamma")

##Most common words in document
tidy(clean_token_dmat) %>%
  filter(document == 6) %>%
  arrange(desc(count))

```

A tibble: 0 x 3

```
## # ... with 3 variables: document <chr>, term <chr>, count <dbl>
## # i Use 'colnames()' to see all variable names
```

```
assignments <- augment(imdb_lda, data = clean_token_dmat)
assignments
```

```
## # A tibble: 5,286 x 4
##   document term      count .topic
##   <chr>    <chr>    <dbl> <dbl>
## 1 35484    timon        25     4
## 2 35484    pumbaa       13     4
## 3 29214    randy        12     1
## 4 12301    christmas    11     6
## 5 320      baseball     10     2
## 6 38255    baseball      3     2
## 7 19086    kite         9      6
## 8 320      burns        8     2
## 9 8473     chris         8     3
## 10 14281    chris         2     3
## # ... with 5,276 more rows
## # i Use 'print(n = ...)' to see more rows
```

##The assignments tibble above count up the words for each topic.

##Term Frequency - Inverse Document Frequency

##Let's look tf-idf to see what is the most important words in the whole reviews.

```
review_tf_idf <- clean_token %>%
  bind_tf_idf(review_index, word, count)
```

##Look at terms with high tf-idf in reviews.

```
review_tf_idf <- review_tf_idf %>%
  arrange(desc(tf_idf))
```

##It looks like the high tf-idf's tf are mostly 1.

##For words that tf=1, it means those words are only contained in one review, and the tf-idf algorithm

##So, let's remove all tf = 1.

```
tf_1 <- which(review_tf_idf$tf==1)
tf_idf_high <- review_tf_idf %>% slice(tf_1) %>% select(-count,-tf,-idf) #remove column 'count','tf','idf'
rm(review_tf_idf)
```

##Now we only have 1993 observations

##Review_index numbers in tf_idf_high

```
index <- unique(tf_idf_high$review_index)
```

##tf-idf plot

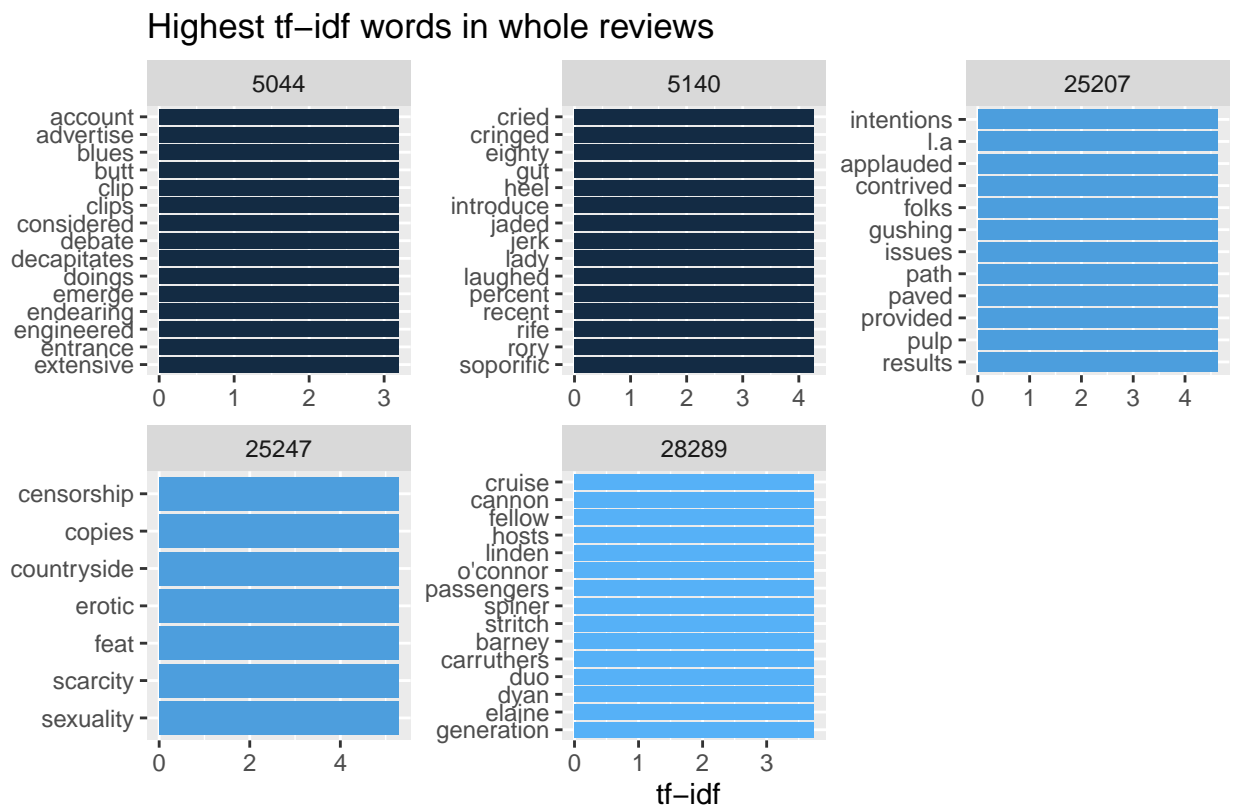
##Let's make the plots with only 6 review_index.

```
tf_idf_high %>%
  filter(review_index %in% c(25207,5044,5140,28289,25247)) %>%
  arrange(desc(tf_idf)) %>%
```

```

group_by(review_index) %>%
distinct(word,review_index, .keep_all = TRUE) %>%
slice_max(tf_idf, n = 15, with_ties = FALSE) %>%
ungroup() %>%
mutate(word = factor(word, levels = rev(unique(word)))) %>%
ggplot(aes(tf_idf, word, fill = review_index)) +
geom_col(show.legend = FALSE) +
facet_wrap(~review_index, ncol = 3, scales = "free") +
labs(title = "Highest tf-idf words in whole reviews",
caption = "IMDB Dataset",
x = "tf-idf", y = NULL)

```



IMDB Dataset

```

##- On each 6 plot, we can see top 15 words with high tf-idf.
##- Among them, we can verify some meaningful words for checking their genres.
##- For example, in review'25247', the words 'censorship','erotic','sexuality'
## imply that the review is about romance movie.

```